

Bayesian ML for Healthcare Applications



Washington
University in St. Louis

SCHOOL OF MEDICINE

Introduction to Bayesian Methods in Healthcare

- Bayesian approach: Incorporates prior knowledge and updates beliefs with new data
- Probabilistic framework: Handles uncertainty in medical data and predictions
- Flexibility: Adapts to various healthcare applications (diagnosis, prognosis, treatment planning)
- Interpretability: Provides probability distributions, not just point estimates
- Scalability: Suitable for both small clinical studies and large-scale health data analysis

Gaussian Processes (GPs) in Healthcare

- Statistical technique for dimensionality reduction of large datasets
- Identifies principal components - directions of maximum variance in the data
- Transforms data into a new coordinate system based on these components
- First principal component captures the most variance, second captures second most, etc.
- Allows data compression by keeping only top N components
- Useful for visualization, noise reduction, and feature extraction in machine learning

GPyTorch for Chest X-ray Classification

- Combines deep learning feature extraction with GP classification
- Feature Extractor: Uses CLIP for efficient image feature extraction
- GP Layer: Implements variational inference for scalable GP classification
- Likelihood: Utilizes SoftmaxLikelihood for multi-class prediction
- Training: Employs VariationalELBO (Evidence Lower Bound) as the objective function
- Uncertainty: Provides prediction probabilities and entropy-based uncertainty estimates



Let's collapse the box

Bayesian Model Components

- Prior: Incorporates domain knowledge about the problem
- Likelihood: Models the data generation process
- Posterior: Updated beliefs after observing the data
- Marginal likelihood: Used for model comparison and hyperparameter tuning
- Predictive distribution: Provides probabilistic predictions for new data points

Advantages of GPs in Medical Imaging

- Uncertainty quantification: Critical for reliable medical decision-making
- Interpretability: Kernel functions can be designed to reflect medical knowledge
- Active learning: Efficiently selects most informative samples for labeling
- Transfer learning: Adapts pre-trained models to specific medical domains
- Robustness: Handles noise and outliers common in medical data

Practical Implementation with GPyTorch

- Integration with PyTorch: Leverages GPU acceleration and automatic differentiation
- Scalability: Uses inducing points for large-scale datasets
- Flexibility: Supports custom mean and kernel functions
- Multi-task learning: Handles multiple related medical tasks simultaneously
- Model selection: Provides tools for kernel and hyperparameter optimization

Limitations of GPs

- **Scalability:** GPs have a computational complexity of making them inefficient for large datasets. Even with GPytorch's optimizations, handling very large datasets remains challenging.
- **Memory Usage:** Storing the covariance matrix requires memory, which can be prohibitive for large datasets, limiting the practical size of the data that can be used.
- **Approximation Quality:** While GPytorch offers methods like variational inference to scale GPs, these approximations may lead to a loss in prediction accuracy, especially in highly non-linear or complex problems.
- **Hyperparameter Sensitivity:** GPs are sensitive to the choice of kernel and its hyperparameters. Poor selection can lead to overfitting or underfitting, requiring careful tuning that can be computationally expensive.
- **Non-Gaussian Noise:** GPs inherently assume Gaussian noise in the data. Dealing with non-Gaussian noise requires custom likelihood functions, which can be complex to implement and computationally expensive.
- **Stationarity Assumption:** Standard GPs assume stationarity in the data, which can be limiting when dealing with non-stationary processes. Extensions to non-stationary kernels are possible but add complexity.
- **Interpretability Challenges:** While GPs provide uncertainty estimates, interpreting these in high-dimensional spaces or complex models can be difficult, making them less intuitive than simpler models.

Evaluation and Interpretation

- Confusion matrix: Assesses model performance across different classes
- Uncertainty visualization: Histograms of prediction entropies
- ROC and PR curves: Evaluates model discrimination ability
- Calibration plots: Ensures reliability of predicted probabilities
- Feature importance: Analyzes contribution of different image regions to predictions

Challenges and Future Directions

- Data scarcity: Developing methods for learning from limited labeled data
- Interpretability: Improving explanations of GP predictions for clinicians
- Computational efficiency: Scaling to larger datasets and real-time applications
- Multi-modal integration: Combining imaging with other clinical data sources
- Causal inference: Incorporating causal relationships in Bayesian models
- Personalized medicine: Adapting models to individual patient characteristics

Summary

- Bayesian methods, particularly GPs, offer powerful tools for healthcare applications
- GPyTorch provides a flexible framework for implementing GP models in medical imaging
- Uncertainty quantification is crucial for responsible AI in healthcare
- Ongoing research aims to address challenges in scalability and interpretability
- The future of Bayesian ML in healthcare promises more accurate, reliable, and personalized medical decision support systems